

# Introduction to the HPC Cluster

Rollins School of Public Health, BIOS 780  
February 2020


# The HPC Cluster

The High Performance Computing (HPC) Cluster is a collection of computers dedicated to (scientific) computation and connected via a dedicated network and storage facility. It runs an operating system called *Linux* (Lee-nuks).

Our Layout:

- Compute nodes
- Submit hosts
- Storage service
- Scheduler
- License Manager





# Submit hosts and Compute Nodes

## Login/Submit hosts:

- Your “login host”. You **ssh**-in to this machine: `ssh <userid>@<submit_host>`
- You most likely will need VPN access to reach these (some campus wired connections are accepted)
- Firewall timeout: 15 minutes
- You copy data onto these systems, which have your home and project directories available (which are mounted everywhere)
- Limited application availability - don't run jobs on these systems
- Named “hpc3”, “hpc4”, “hpc5”, etc.

## Compute nodes:

- These are where the jobs are run
- Applications readily available here
- No direct login access
- Many cores, large amounts of RAM
- Physical servers (not VMs)
- Named after fruit (“apple”, “banana”, “cherry”) or research group (“gene1”, “gene2”, etc.) or randomly (“dynareg1”)
- There are 32 compute nodes
- 664 cores -> 1328 threads total
- Running RHEL 7

# Logging in to the Cluster

- We use the Secure Shell (**ssh**) program to log into the submit hosts
- The easiest operating systems: Linux and Mac OS X clients. Use the provided **Terminal** program, and **ssh** is already included.



# Preliminaries: Getting an SSH client for Windows

- Connecting to the cluster from Windows is more complicated. You'll need to add/install an **ssh** client. There are at least three options:
  - Install **Putty**. Go to [putty.org](https://putty.org), download 64-bit .msi file and install. It works on older Windows systems
  - On Windows 10, use **Windows Subsystem for Linux**. Download Ubuntu from Windows store
  - Install **Linux** via a virtual machine: e.g., VirtualBox or VMWare.
  - Excellent how-to video, *Learning Linux Command Line* at
  - <https://www.linkedin.com/learning-login/> (use Emory email)
  - HelpDesk: installation of Putty

# Preliminaries: Changing your unix \$SHELL

- A legacy “feature” of central IT at Emory: everyone’s Unix shell is set to ‘ksh’ in LDAP directory service.
- No one uses **ksh** anymore (the default shell on Linux is **bash**, Mac OS X now uses **zsh**). The **bash** shell is preferred and supported on the cluster.
- All new Emory users have to change it themselves:  
<https://mynetid.emory.edu/IDMProv/portal/cn/DefaultContainerPage/WelcometoMyNetID>
- Call 7-7777 (Central IT help line) for assistance.



# More Preliminaries: VPN access

- If you connect to the cluster from an off-site location, including Emory Unplugged (Emory's wireless network), you'll need access to the Emory VPN.
- This is granted by individual request by central IT (*LITS*):  
<https://it.emory.edu/security/vpn.html>
- You shouldn't need to request anything other than having VPN access "activated" for your Emory NetID.
- You need to install Emory's VPN client on your computer.
- Call 7-7777 (Central IT help line) for assistance.





# Using SSH

- Use command `ssh <username>@<submit_host>.sph.emory.edu`, e.g., “`ssh fred@hpc6.sph.emory.edu`”
- Accept new key the first time you login
- Best practice: password-less key exchange and `ssh-agent`
- **ssh-copy-id** works well from any Linux/Mac OS X terminal.
- **.ssh/config** can be used to save aliases to further reduce typing.



# Staying Connected to the Cluster: **screen**

- Once logged in, both the VPN and the network firewall settings can “time out”.
- As a work-around, use the Unix command **screen**.
- After logging in, simply type ‘screen’. It will return a prompt to you, as if nothing happened. Type “<ctrl>-a d”, to exit a screen.
- If your connection gets dropped (while in a screen), you can “reattach” to the session via **screen -r**. (Type ‘<ctrl>-a ?’ in **screen**, for command list.)



# Unix: a few helpful commands

- **du** : shows disk usage. Ex: **du -hs /home/<netid>** shows how large your home directory is (in gigabytes). (*This command may take a few minutes to complete.*)
- **ls** : list files in current working directory
- **pwd** : print working directory
- **more** : read contents of a text file, with paging
- **chown** : change user or group ownership of a file
- **vim** : a lightweight text file editor (worth learning)
- **man** : read manual pages. Ex. **man chown**
- **scp** : secure copy, part of **ssh**
- Free course! *Learning Linux Command Line* via Emory
- Go to <https://www.linkedin.com/learning-login/>

# Using the Cluster: Grid Engine q-commands

Work on the cluster is submitted in the form of “jobs”, or programming tasks, that are sent to the compute nodes from the submit nodes. The *q-commands* are Grid Engine commands that perform various job functions:

Submit jobs	<b>qsub, qlogin</b>
-------------	---------------------

Check job status	<b>qstat</b>
------------------	--------------

Delete jobs	<b>qdel</b>
-------------	-------------

Cluster/queue state	<b>qhost, qstat</b>
---------------------	---------------------



# Cluster status: **qhost** - show available hosts

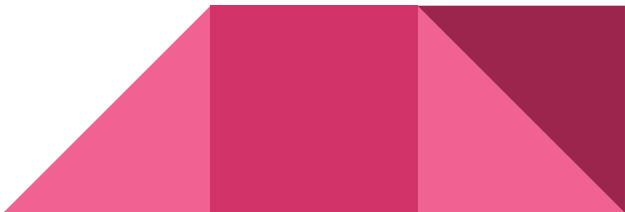
```
khaynes@hpc1 ~->qhost
```

HOSTNAME	ARCH	NCPU	NSOC	NCOR	NTHR	NLOAD	MEMTOT	MEMUSE	SWAPTO	SWAPUS
global	-	-	-	-	-	-	-	-	-	-
alpha	lx-amd64	72	2	36	72	0.04	220.3G	5.8G	4.0G	1.4G
alz1	lx-amd64	56	2	28	56	0.00	251.8G	5.0G	128.0G	595.5M
apple	lx-amd64	40	2	20	40	0.01	251.9G	2.8G	128.0G	2.5G
banana	lx-amd64	40	2	20	40	0.01	251.9G	3.0G	128.0G	2.6G
bravo	lx-amd64	24	2	12	24	0.08	62.9G	5.5G	1024.0M	1023.7M
cbis1	lx-amd64	40	2	20	40	0.03	125.7G	2.4G	4.0G	1.1G
cbis2	lx-amd64	40	2	20	40	0.00	125.7G	2.5G	4.0G	1.2G
charlie	lx-amd64	24	2	12	24	0.04	62.9G	1.4G	13.6G	1.2G
cherry	lx-amd64	32	2	16	32	0.06	125.9G	3.4G	128.0G	654.4M
condor	lx-amd64	40	2	20	40	0.09	125.8G	59.7G	96.0G	176.7M
delta	lx-amd64	24	2	12	24	0.13	62.9G	6.1G	13.6G	1.1G
dynareg1	lx-amd64	48	2	24	48	1.01	251.6G	14.6G	4.0G	264.2M
eagle	lx-amd64	8	2	4	8	0.17	31.4G	1.1G	64.0G	924.7M
echo	lx-amd64	24	2	12	24	0.13	62.9G	13.5G	64.0G	1.5G
foxtrot	lx-amd64	32	2	16	32	0.07	125.9G	2.9G	128.0G	1.1G
gene1	lx-amd64	40	2	20	40	0.02	125.8G	3.4G	128.0G	1.0G
gene10	lx-amd64	72	2	36	72	0.00	251.8G	4.3G	128.0G	485.8M

# Cluster status: qstat -f

```
khaynes@hpc1 ~->qstat -f -q long.q
queue name      qtype resv/used/tot. np_load  arch      states
-----
long.q@alpha    BP    0/2/72        0.04    lx-amd64
long.q@bravo    BP    0/2/24        0.08    lx-amd64
long.q@charlie  BP    0/1/24        0.04    lx-amd64
long.q@delta    BP    0/2/24        0.13    lx-amd64
long.q@echo     BP    0/3/24        0.13    lx-amd64
long.q@foxtrot  BP    0/1/32        0.07    lx-amd64
long.q@hotel    BP    0/1/32        0.05    lx-amd64
long.q@india.sph.emory.edu BP    0/1/48        0.02    lx-amd64
long.q@juliett.sph.emory.edu BP    0/2/48        0.04    lx-amd64
long.q@kilo.sph.emory.edu BP    0/2/48        0.04    lx-amd64
```

# Submitting jobs: **qsub**

- Batch job submission is done via **qsub**
  - **qsub <program\_name>** submits a job running <program\_name> with default attributes, after which the job will appear in the output of **qstat**
  - Many attributes can be specified, such as run queue, input, output, job name, memory usage, etc.
  - Ex. **qsub -cwd -i /data/example.in -o /results/example.out -N testrun1 -q long.q -l mem\_free=2GB R\_wrapper\_script.sh**
  - That's a lot of typing! Most people put those attributes in a job submission script.
- 

File Edit View Search Terminal Help

```
#!/bin/bash
```

```
# set the name
```

```
#$ -N testrun1
```

```
# Set the input file
```

```
#$ -i /data/example.in
```

```
# Set the output, both stdout and stderr
```

```
#$ -o /results/example.out -o /results/example.err
```

```
# other submit options
```

```
#$ -l mem_free=2GB
```

```
#$ -q long.q
```

```
some_R_program.R
```

```
# END
```

A sample job submission script.



```

File Edit View Search Terminal Help
guest@hpc5~> cat GE_TEST/bin/simple.sh
#!/bin/bash

# This is a simple example of a SGE batch script

# request Bourne shell as shell for job
#$ -S /bin/bash

# print date and time
date
# Sleep for 20 seconds
sleep 20
# print date and time again
date

guest@hpc5~> qsub GE_TEST/bin/simple.sh
Your job 1606884 ("simple.sh") has been submitted
guest@hpc5 ~>qstat
job-ID      prior    name             user          state submit/start at   queue                          jclass      slots ja-task-ID
-----
 1606884    0.55476 simple.sh        guest         r       02/07/2019 15:34:02 short.q@gene10.sph.emory.edu          1

guest@hpc5~> ls -l simple.sh.*
-rw-r--r-- 1 guest RSPH  0 Feb  7 15:34 simple.sh.e1606884
-rw-r--r-- 1 guest RSPH 29 Feb  7 15:34 simple.sh.o1606884

guest@hpc5~> more simple.sh.o1606884
Thu Feb  7 15:34:02 EST 2019
Thu Feb  7 15:34:22 EST 2019

```

A **qsub** workflow example.

# Checking on Jobs: **qstat** and **qacct**

- **qstat** (without arguments) shows the status all of your currently queued jobs
- Jobs usually have status of **r** (“running”), **qw** (“queue wait” or pending), or **Eqw** (in and “error” state)
- Use **qacct -j <job\_id>** to get detailed information about any job.



# Killing jobs: **qdel**

- Use **qdel** *<job\_id>* to kill a submitted job



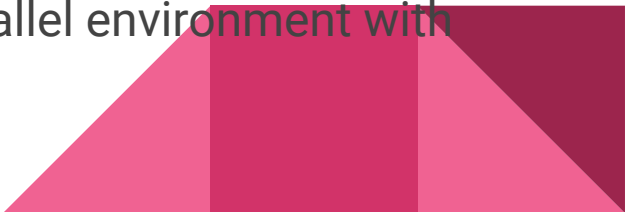
```

File Edit View Search Terminal Help
user@hpc5> ./sleeper.sh 30 1
Here I am. Sleeping now at: Thu Feb 7 16:06:48 EST 2019
Now it is: Thu Feb 7 16:07:18 EST 2019
user@hpc5> qsub sleeper.sh 30 1
Your job 1606886 ("Sleeper") has been submitted
user@hpc5> qsub sleeper.sh 30 1
Your job 1606887 ("Sleeper") has been submitted
user@hpc5> qsub sleeper.sh 30 1
Your job 1606888 ("Sleeper") has been submitted
user@hpc5> qstat
job-ID      prior    name         user          state submit/start at     queue                          jclass      slots ja-task-ID
-----
    1606886  0.55476  Sleeper      user          r      02/07/2019 16:07:31 short.q@gene9.sph.emory.edu          1
    1606887  0.55476  Sleeper      user          r      02/07/2019 16:07:35 short.q@gene10.sph.emory.edu         1
    1606888  0.55476  Sleeper      user          r      02/07/2019 16:07:35 short.q@gene3                      1
user@hpc5> qdel 1606887
user has registered the job 1606887 for deletion
user@hpc5> qstat
job-ID      prior    name         user          state submit/start at     queue                          jclass      slots ja-task-ID
-----
    1606886  0.55476  Sleeper      user          r      02/07/2019 16:07:31 short.q@gene9.sph.emory.edu          1
    1606887  0.55476  Sleeper      user          dr     02/07/2019 16:07:35 short.q@gene10.sph.emory.edu         1
    1606888  0.55476  Sleeper      user          r      02/07/2019 16:07:35 short.q@gene3                      1
user@hpc5> qstat
job-ID      prior    name         user          state submit/start at     queue                          jclass      slots ja-task-ID
-----
    1606886  0.55476  Sleeper      user          r      02/07/2019 16:07:31 short.q@gene9.sph.emory.edu          1
    1606888  0.55476  Sleeper      user          r      02/07/2019 16:07:35 short.q@gene3                      1
khaynes@hpc1 ~>

```

A **qdel** workflow example.

# Cluster queues

- *qlogin.q* - interactive sessions; dedicated 64 cores; timeout after 24 hours
  - *long.q* - general use; no time limit; 14 jobs/user max.
  - *short.q* - general use subordinate queue; 72 hour limit; 240 cores; limited to 40 jobs/user
  - Restricted queues: *fruit.q*, *gene.q*, *sunlab.q*, *cbis.q*, *alz.q*, *mh.q*
  - Restricted queues have priority on owned systems, and have no per user job limit or run-time limit
  - All queues have corresponding parallel queue (*long*, *short*, *gene*, etc. )
  - Ex. **qsub -pe long 4 <program\_name>** requests a parallel environment with four processors
- 

# Available Software

- C/C++, Fortran compilers
- Python\*
- R\*\*
- MATLAB
- SAS
- Openmpi
- Others, including user-installed

Look in /usr/local/ for the latest installed versions!

( E.g., use /usr/local/R-3.6.0/bin/R)




# Storage on Cluster

- Each account is granted a home directory in /home/<NetID>, which is mounted across all nodes of the cluster (i.e., software installed there is available everywhere).
- Initial quota is set to 16 GB. Quota can be increased once, no questions asked. Users are responsible for deleting/off-loading and compressing files.
- For large data installations, request a project directory. There is a cost associated with project directories, depending on type.



# Getting Help

- If you need help with the HPC Cluster, you should open a help ticket with RSPH IT.
  - Send email to “[help@sph.emory.edu](mailto:help@sph.emory.edu)” and CC: [keven.haynes@emory.edu](mailto:keven.haynes@emory.edu), and include the following:
    - Your Emory NetID
    - **State in both the Subject and the body of the ticket that you your request concerns the HPC Cluster.**
    - The particular issue you are having. (“It is broken” is not an issue.)
    - The name of the system you were logged into, or the particular software package.
  - We can help with systems-related issues.
- 



# More info

- Excellent how-to videos are available to Emory community (free!) via <https://www.linkedin.com/learning-login/>. Click on sign in, then enter your Emory email address, and you will be redirected to Emory for sign in. Search for *Learning Linux Command Line* for video instructions on how to install Windows Subsystem for Linux or Linux in a virtual machine (plus, basic Linux usage).
- The old HPC Cluster documentation is still here:  
<https://intranet.sph.emory.edu/services/it/environment/cluster/index.html>

(It is in need of updating!)

